

Chapter 2

RDF Syntax 2



Topics

- Basic concepts of RDF
 - Resources, properties, values, statements, triples
 - URIs and URIrefs
 - RDF graphs
 - Literals, qnames
- Vocabularies and modeling
 - Vocabularies
 - Blank nodes, data modeling, types, reification
 - Lists, bags, collections
- Serialization of RDF graphs
 - XML, Turtle, Ntriples
- Critique of RDF

Types

RDF type

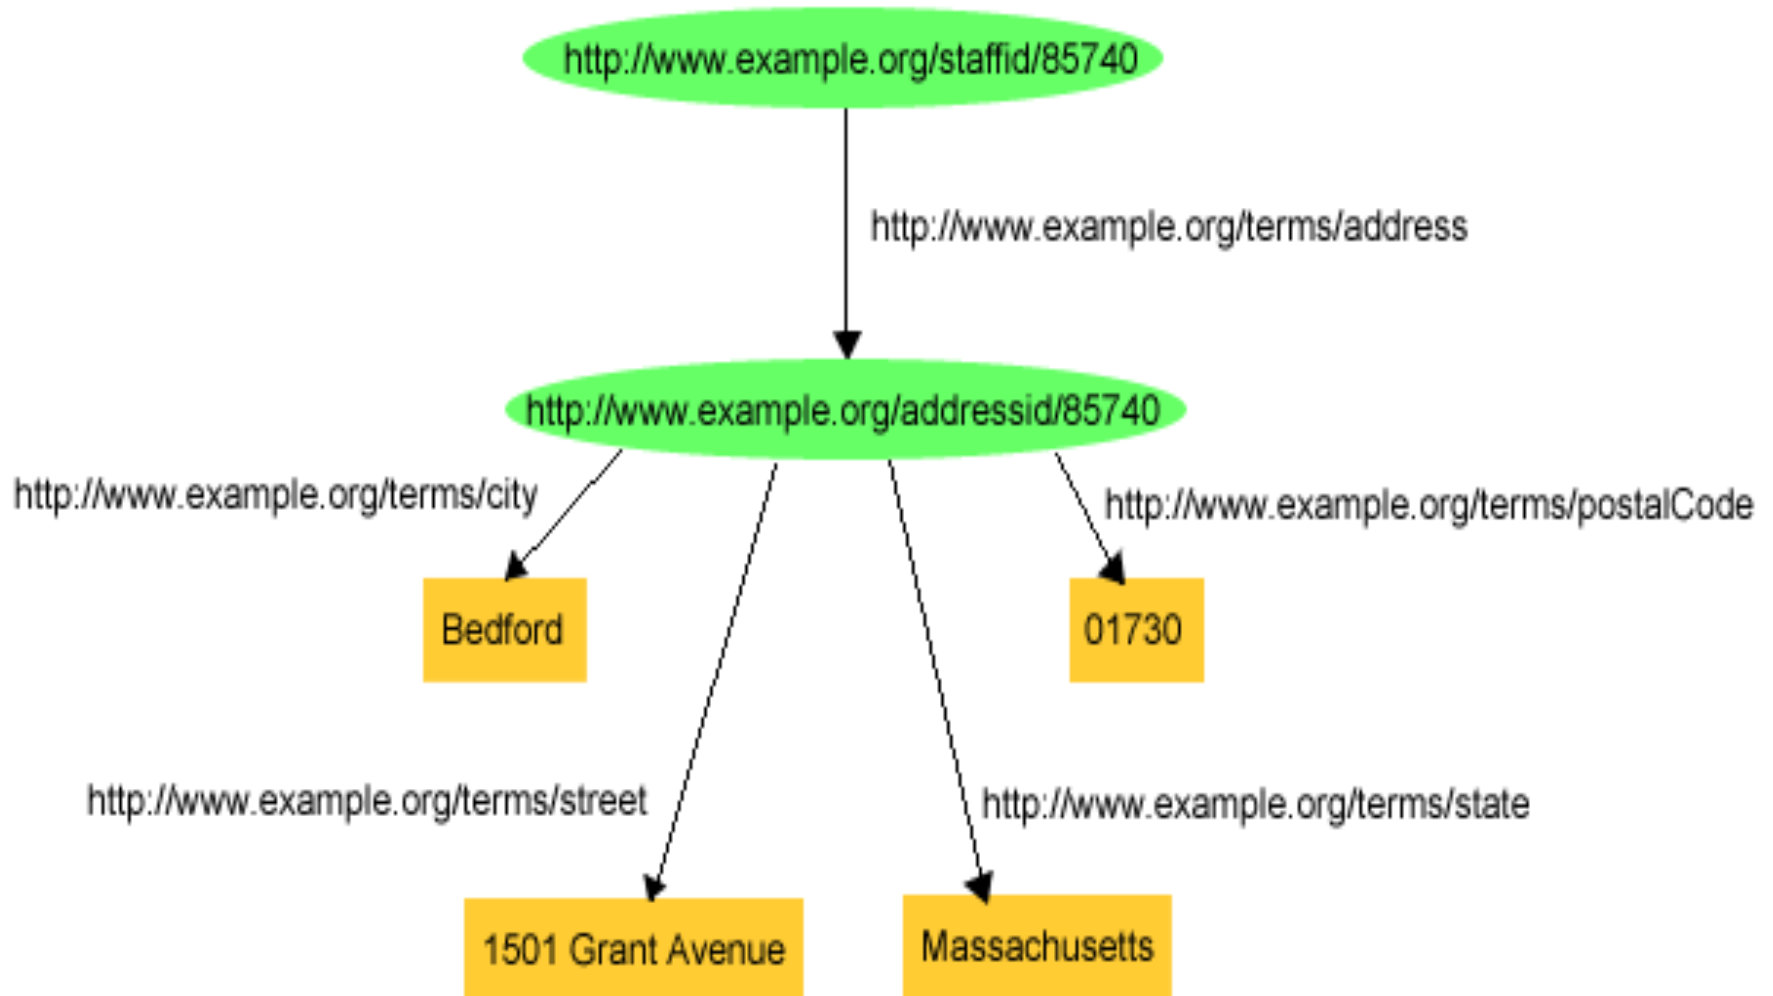
- RDF has a type predicate that links a resource to another that denotes its type
 - `ex:john rdf:type foaf:Person .`
 - `<http://example.org/john <http://www.w3.org/1999/02/22-rdf-syntax-ns#type <http://xmlns.com/foaf/0.1/Person> .`
- RDFS adds sub-type concept & constraints between predicates & types of their arguments
- OWL adds still more concepts operating on types

Data Modeling

Structured Values in RDF

- Given the triple like:
ex:857 exstaff:address "15 Grant Ave, Bedford, MA 01730".
- How can we best represent separate information for the street, city, state and zip code?
- Two possibilities:
 - Use four predicates (e.g., exstaff:street_address, ...) to associate values with exstaff:857
 - Create an address resource to attach the four predicates to and link that to exstaff:address with the ex:address predicate

Structured Values in RDF



Structured Values in RDF

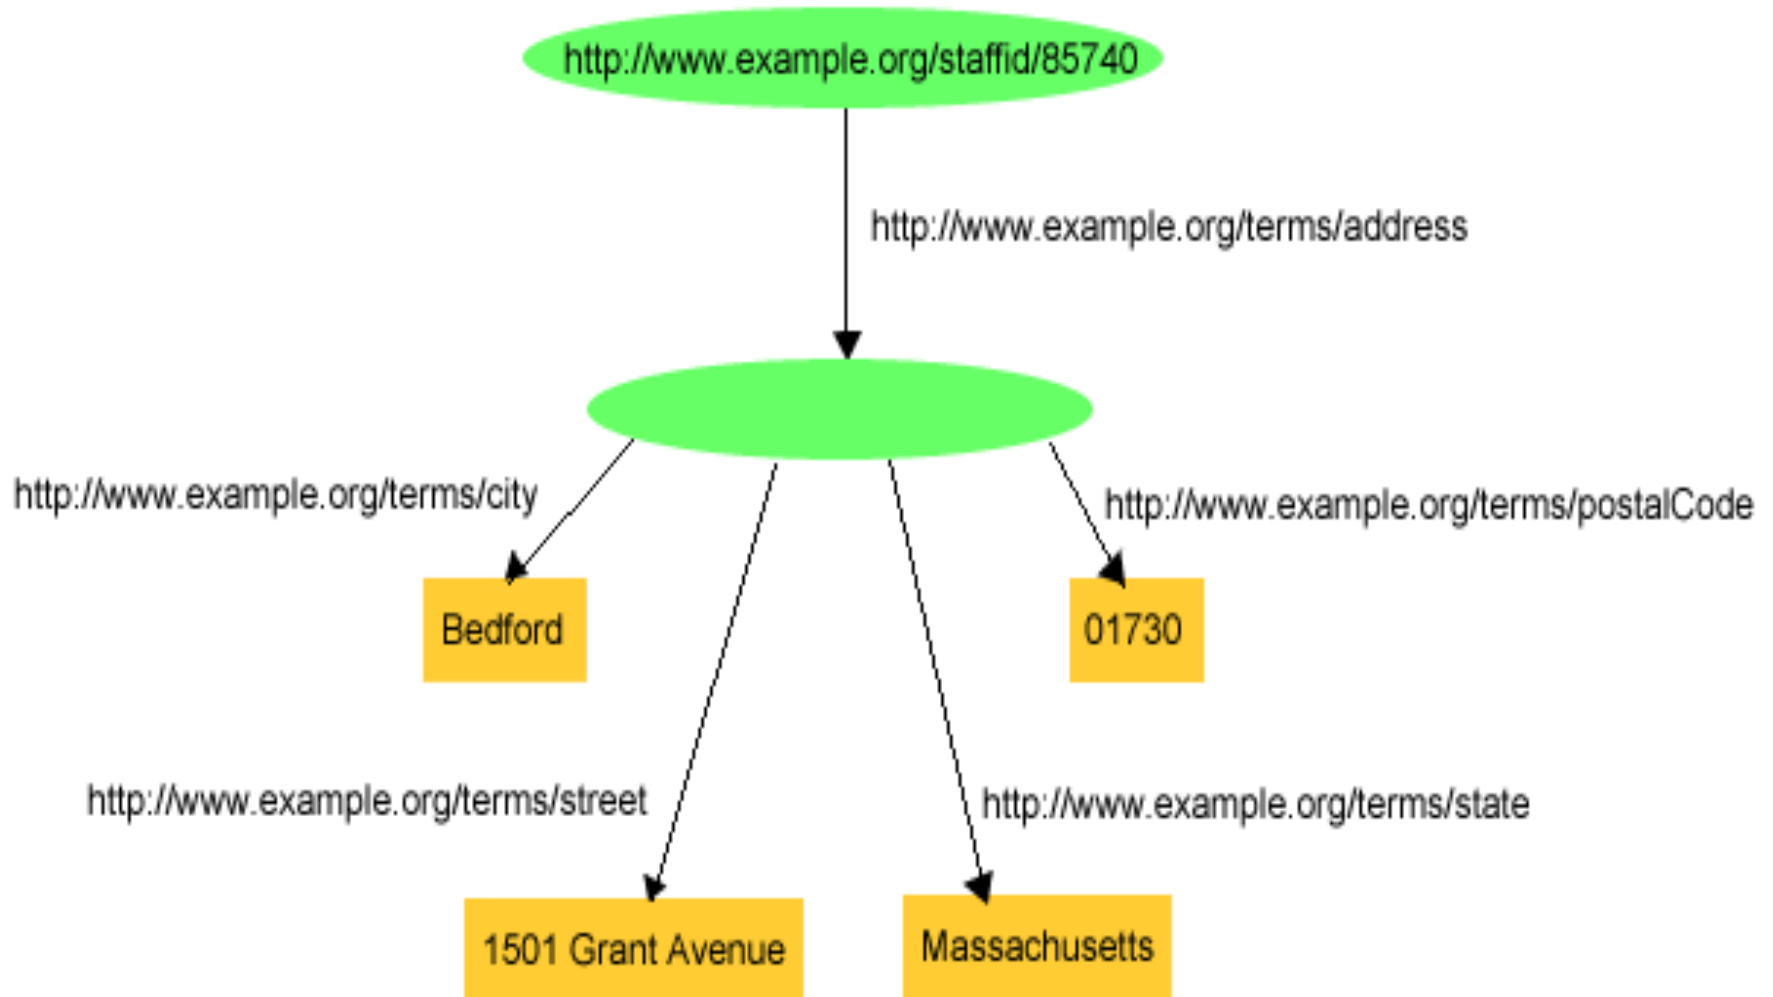
Pr as triples:

```
exstaff:85740 exterms:address exaddressid:85740 .  
exaddressid:85740 exterms:street "1501 Grant Ave" .  
exaddressid:85740 exterms:city "Bedford" .  
exaddressid:85740 exterms:state "MD" .  
exaddressid:85740 exterms:postalCode "01730" .
```


Structured Values in RDF

- This approach involves adding **many “inter-mediate” URIs** (e.g., exaddressid:85740) for aggregate concepts like John's address
- Such concepts may never need to be referred to directly from outside a particular graph, and hence **may not require “universal” identifiers**
- RDF allows us to use **blank nodes** and **blank node identifiers** to deal with this issue
 - Node IDs in the `_` namespace are bnodes, e.g. `_:`

Blank Node, aka bnode



Blank Nodes Using Triples

exstaff:85740 exterms:address ?? .

?? exterms:postalCode "01730" .

Exstaff:72120 exterms:address ??? .

??? exterms:postalCode "01702" .

- We want to ensure that the bnodes for 85740's and 72120's addresses are distinct
- The graphical notation does this by using two different objects for the bnodes
- RDF allows us to assign an special ID to a bnode while still maintaining its blank node nature


Blank Node Identifiers

```
exstaff:85740 exterm:address _:johnaddress .  
_:johnaddress exterm:street "1501 Grant Avenue" .  
_:johnaddress exterm:postalCode "01730" .
```

- Distinct bnode must have **different** bnode ids
- Bnode ids have significance only in a **single** graph
 - *dbpedia:Alan_Turing* refers to the same thing in every graph, but a bnode `_:1` in two different graphs may not
 - Merging two graphs requires us to rename their bnode ids to avoid accidental conflation (e.g., `_:1 => _:100`)
- Bnode ids may only appear as subjects or objects and **not as predicates** in triples

Semantics of Blank Nodes

- In terms of **first-order logic**, blank nodes correspond to existentially quantified variables
- Another example: “John’s mother is 50”
- **FOL**: $\exists x \text{ mother}(\text{john}, x) \wedge \text{age}(x, 50)$
- **RDF**: `:john :mother _32 . :_32 :age "50" .`



Blank nodes are good for

- Representing **n-ary relationships** in RDF
e.g., the relationship between John Smith and the street, city, state, and postal code components of his address
- To make statements about **resources that don't have URIs** but are described by relationships with other resources that do
e.g., John's mother

Example

- To make statements about Jane Smith we could use her email address URI (<mailto:jane@example.org>) to denote her
- Well, if we do so, how are we going to record information both about **Jane's mailbox** (e.g., the server it is on) as well as about **Jane herself** (e.g., her current physical address)? Similarly, if we use her Web page URI etc.

Bnode Example

When Jane herself does not have a URI, a blank node provides a better way of modeling this situation

```
_:jane exterm:mailbox <mailto:jane@example.org> .  
_:jane rdf:type exterm:Person .  
_:jane exterm:name "Jane Smith" .  
_:jane exterm:emplID "23748" .  
_:jane exterm:age "26" .
```


Another use case: Measurements

- What does this mean?

[dbr:Nile](#) [dbp:length](#) "6853"^^xsd:integer

- Click on [dbp:length](#) to see its definition

The screenshot shows a web browser window with the address bar displaying "dbpedia.org/page/Nile". The page title is "About: Nile". Below the title, it states "An Entity of Type : [body of water](#), from Named Graph : <http://dbpedia.org>, within Data Space : [dbpedia.org](#)". The main text describes the Nile as a major north-flowing river in northeastern Africa, 6,853 km long, and an "international" river. Below the text is a table with two columns: "Property" and "Value".

Property	Value
dbo:Stream/discharge	▪ 2830.0
dbo:abstract	▪ The Nile (Arabic: النيل, Eg. en-Nīl, Std. an-Nīl; Coptic: ⲡⲓⲁⲣⲱ, P(h)iario; Ancient Egyptian: Ḥ'pī and Iteru) is a major north-flowing river in northeastern Africa,

Another use case: Measurements

- What does this mean?

[dbr:Nile](#) [dbp:length](#) "6853"^^xsd:integer

- We can click on [dbp:length](#) to see its definition

[dbp:length](#) rdf:type rdf:Property .

[dbp:length](#) rdfs:label "Length"@en .

- Unfortunately, the definition doesn't specify the unit of measurement 😞

Another use case: Measurements

- What does this mean?

dbr:Nile dbp:length "6853"^^xsd:integer

- Measurements typically have a numeric *value* and a *unit*
 - **Weight:** 2.4 pounds vs. 2.4 kilograms
 - **Length:** 5 miles vs. 5 kilometers
 - **Price:** 29.00 in US Dollars vs. 21.16 Euro
 - **Time:** 30 years vs. 3 milliseconds
- We can use a bnode to represent a measurement as a pair with a value and unit

Measurements


- What does this mean?

dbr:Nile dbp:length _:1 .

_:1 rdf:type ex:Measure .

_:1 rdf:value "6853"^^xsd:integer .

_:1 un:units dbr:Kilometre .



- The RDF namespace has a *value* property but assigns no specific meaning to it

Serialization

RDF Serialization

- Abstract model for RDF is a graph
- Serialize as text for exchange, storage, viewing and editing in text editors
- The big three
 - XML/RDF – the original
 - Ntriples – simple, but verbose; good for processing
 - Turtle – compact, easy for people to read and write
- Special formats
 - Trig – a format for named graphs
 - RDFa – embed RDF in HTML attributes
 - JSON-LD – RDF statements as a JSON object

XML encoding for RDF

```
<rdf:RDF xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
  xmlns:dc="http://purl.org/dc/elements/1.1/"
  xmlns:bib="http://daml.umbc.edu/ontologies/bib/">
<rdf:Description about="http://umbc.edu/~finin/talks/idm02/">
  <dc:title>Intelligent Information Systems on the Web </dc:title>
  <dc:creator>
    <rdf:Description >
      <bib:name>Tim Finin</bib:name>
      <bib:email>finin@umbc.edu</bib:email>
      <bib:aff resource="http://umbc.edu/" />
    </rdf:Description>
  </dc:creator>
</rdf:description>
</rdf:RDF>
```

RDF/XML is a W3C
Standard widely used for
storage and exchange

Being supplanted by other
forms

Complex and confusing so we
won't spend time on it

Ntriples

- Good for ingesting into a program or store
- Sequence of triples each terminated with a “.”
- URIs encased in angle brackets; no QNames; literals in double quotes
- Trivial to parse/generate; common download format for RDF datasets (e.g., [DBpedia](#))
- Uses lots of characters due to repeated URLs, but compresses well

[W3C Specification](#)

```
<http://example.org/Turing><http://www.w3.org/1999/02/22-rdf-syntax-ns#type> <http://xmlns.com/foaf/0.1/Person> .  
<http://example.org/Turing> <http://xmlns.com/foaf/0.1/name> "Alan Turing" .  
<http://www.w3.org/2001/sw/RDFCore/ntriples/> <http://xmlns.com/foaf/0.1/mbox> <mailto:alan@turing.org> .
```


Turtle

- Ntriples \subset Turtle \subset N3
- Compact, easy to read and write and parse
- Qnames, [] notation for blank nodes, ; and ,

@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .

@prefix dc: <http://purl.org/dc/elements/1.1/> .

@prefix foaf: <http://xmlns.com/foaf/0.1/> .

<http://www.w3.org/TR/rdf-syntax-grammar>

dc:title "RDF/XML Syntax Specification (Revised)" ;

dc:creator [foaf:name "Dave Beckett";

foaf:mbox <mailto:dave@beckett.org> ,

<mailto:dbeck@gmail.com>

].

Some details

- @PREFIX lines define namespace abbreviations
- Basic pattern is
Subj pred1 value1;
pred2 value2;
pred3 value3, value4 .
- Special notation for the rdf:type predicate
:john a **foaf:Person**; foaf:name "John Smith" .
- Special notation for anonymous bnodes
:john foaf:knows [a foaf:Person; foaf:nick "Bob"].

Notation3 or N3

- N3 was an early turtle-like notation developed by Sir Tim_Berners Lee himself
- Included support for inference rules
 - See [CWM](#) for software
- Never became a recommended W3C standard
 - Some of its features were problematic for OWL
 - Supplanted by Turtle

Try...

- Some simple RDF serialization examples
- Simple.ttl

A simple Turtle example

@prefix foaf: <http://xmlns.com/foaf/0.1/> .

@prefix : <#> .

:john a foaf:Person;

foaf:gender "Male";

foaf:name "John Smith", "Johnny Smith";

foaf:knows :mary,

[a foaf:Person;

foaf:mbox <mailto:mary.smith@gmail.com>] .

:mary a foaf:Person;

foaf:name "Mary Smith" .

Notation translation

- Most modern Semantic Web software can read and write rdf in all major serializations
 - E.g., Protégé, Jena, Sesame,
- There are also simple programs that can convert between them
 - [rdf2rdf](#) is an example written in Java

Reification

Reification

- Sometimes we wish to make **statements about other statements**

E.g., to record provenance data, probability, or to assert
:john :believes { :mary :loves :john }

- We must be able to refer to a statement using an identifier
- RDF allows such reference through a reification mechanism which turns a statement into a resource

Reify

- Etymology: Latin *res* thing
- Date: 1854
- To regard (something abstract) as a material or concrete thing

Wikipedia: reification (computer science)

Reification is the act of making an abstract concept or low-level implementation detail of a programming language accessible to the programmer, often as a first-class object. For example,

- The C programming language reifies the low-level detail of memory addresses
- The Scheme programming language reifies continuations (approximately, the call stack)
- In C#, reification is used to make parametric polymorphism implemented as generics a first-class feature of the language
- ...

Reification Example

:949352 uni:name "Grigoris Antoniou" .

reifies as

```
[a rdf:Statement;  
  rdf:subject: :949352  
  rdf:predicate uni:name;  
  rdf:object "Grigoris Antoniou" ] .
```

Another reification example

“Alice suspects that Bob loves Carol”

@prefix ep: <<http://example.com/epistemology>>

@prefix rdf: <<http://www.w3.org/1999/02/22-rdf-syntax-ns#>>.

@prefix xsd: <http://www.w3.org/2001/XMLSchema>

:bob :loves :carol .

[:alice ep:believes

 [a rdf:Statement;

 rdf:subject :bob;

 rdf:predicate :loves;

 rdf:object :carol;

 ex:certainty “0.50”^^xsd:integer]

Reification

- **rdf:subject**, **rdf:predicate** & **rdf:object** allow us to access the parts of a statement
- The **ID** of the statement can be used to refer to it, as can be done for any description
- We write an **rdf:Description** if we don't want to talk about a statement further
- We write an **rdf:Statement** if we wish to refer to a statement

Containers

Container Elements

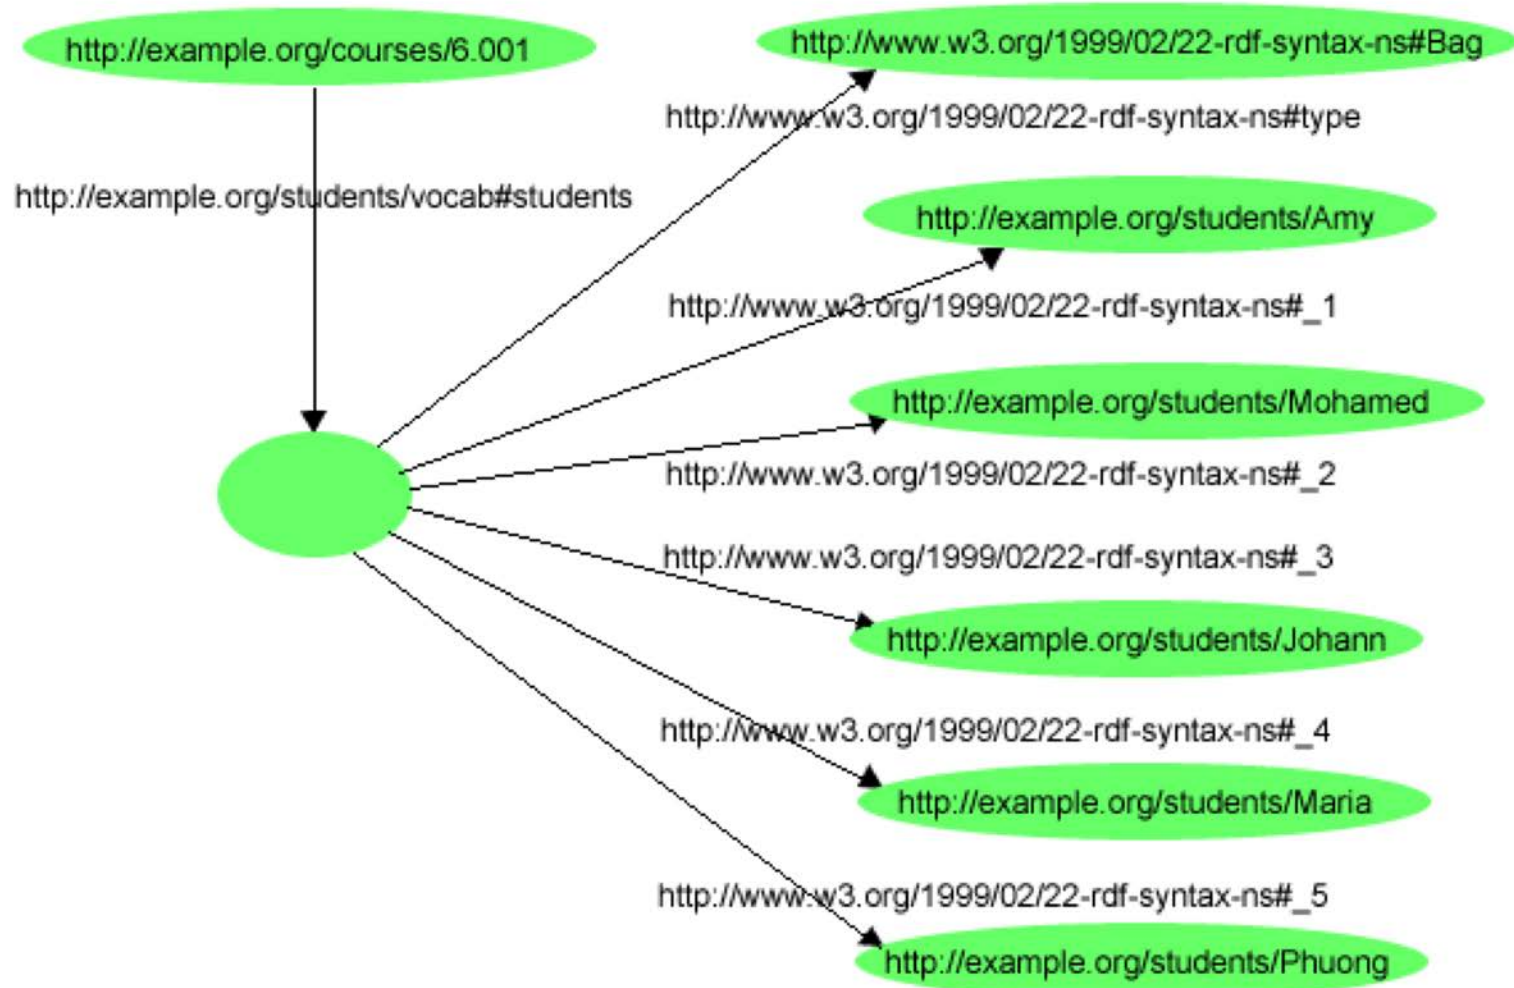
- RDF has some vocabulary to describe collections of things and make statements about them
- E.g., we may wish to talk about the courses given by a particular lecturer
- The content of container elements are named **rdf:_1**, **rdf:_2**, etc.
 - Alternatively **rdf:li**
- Containers seem a bit messy in RDF, but are needed
- :john :teaches [a rdf:Bag; rdf:li :cmisc201, :cmisc202, cmisc345 .] .

Three Types of Container Elements

- **rdf:Bag** an unordered container, allowing multiple occurrences
e.g., members of the faculty, documents in a folder
- **rdf:Seq** an ordered container, which may contain multiple occurrences
e.g., modules of a course, items on an agenda, alphabetized list of staff members
- **rdf:Alt** a set of alternatives
e.g., the document home site and its mirrors, translations of a document in various languages

Example for a Bag

Let's describe a course with a collection of students



Example for a Bag

@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>.

@prefix s: <http://example.org/students/vocab#>.

<http://example.org/courses/6.001>

s:students [

a rdf:Bag;

rdf:_1 <http://example.org/students/Amy>;

rdf:_2 <http://example.org/students/Mohamed>;

rdf:_3 <http://example.org/students/Johann>;

rdf:_4 <http://example.org/students/Maria>;

rdf:_5 <http://example.org/students/Phuong>.

].

Bags and Seqs are never full!

- RDF's semantics is "open world", so...
 - Not possible "to close" the container, to say: "these are **all** elements, there are no more"
 - RDF is a graph, with no way to exclude the possibility that there is another graph somewhere describing additional members
- Lists are collections with only the specified members mentioned.
- Described using a linked list pattern via:
 - `rdf:List`, `rdf:first`, `rdf:rest`, `rdf:nil`

Open vs. closed world semantics

- Reasoning systems make a distinction between open and closed world semantics
 - OWS: being unable to prove that something is true or false says nothing about its veracity
 - CWS: what cannot be proven to be true is false
- Default model for Semantic Web is OWS
 - This was a design decision made early on

Open vs. closed world semantics

- Classical logic uses Open World Semantics

Being unable to prove $P=NP$ doesn't convince us that it's false

- Database systems typically assume CWS

The DB includes all trains between NYC and DC

- Prolog's unprovable operator (not or $\backslash +$) supports CWS

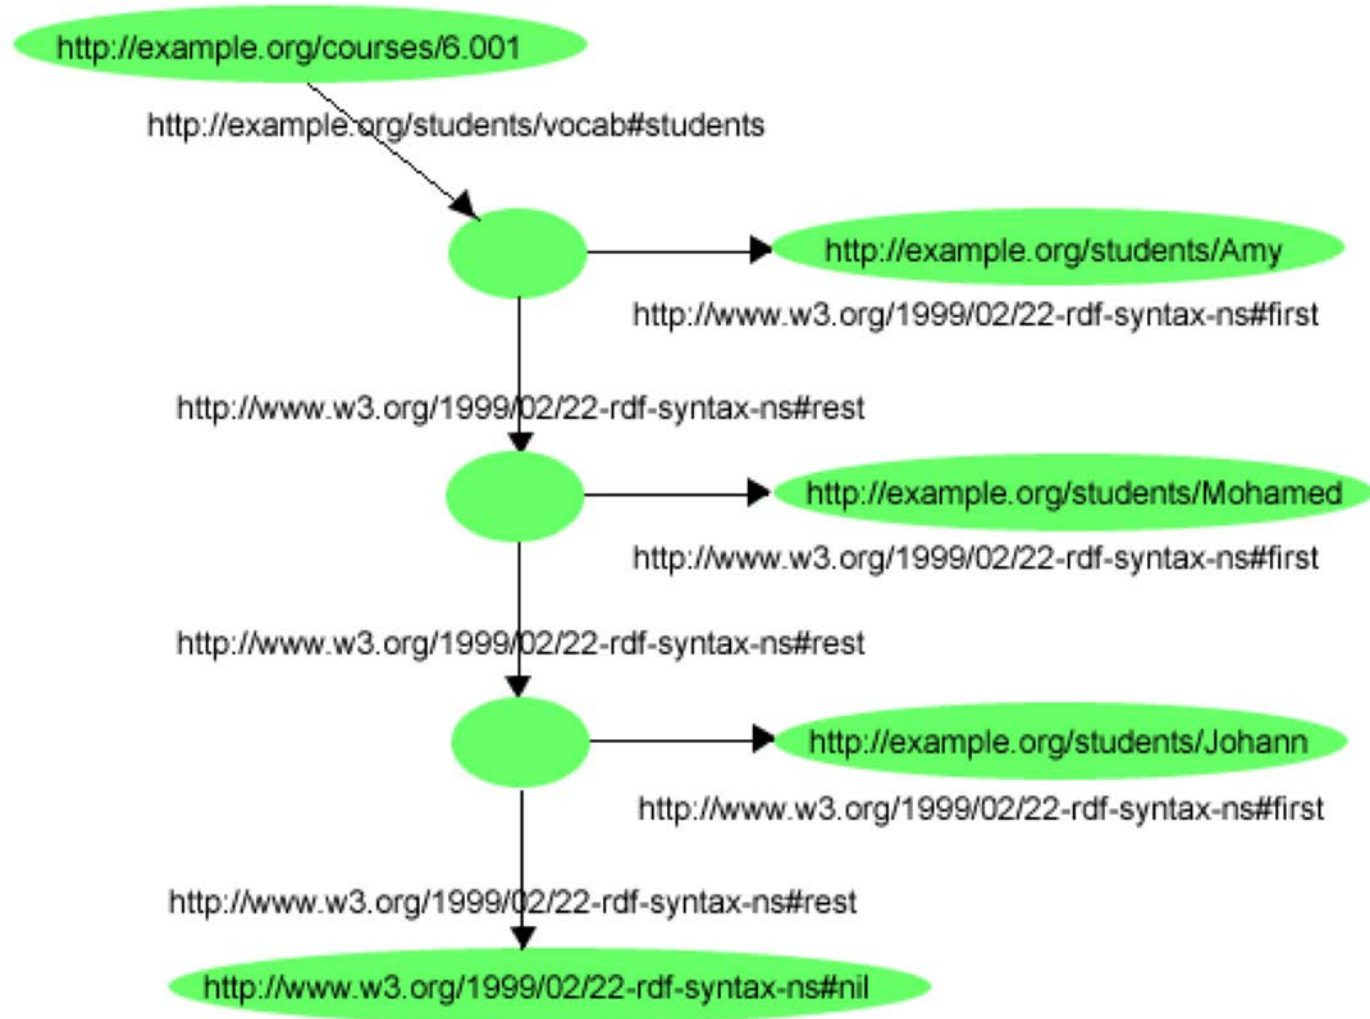
`flys(x) :- bird(x), $\backslash +$ flightless(x).`

`flightless(x) :- penguin(x); ostrich(x); emu(x).`

- Some systems let us specify for which predicates we have complete knowledge and for which we don't
 - If UMBC's DB doesn't list you as registered for CMSC691, you are not registered
 - UMBC's DB system knows some of your minors but not all

RDF Lists

An ordered list of the three students in a class



RDF Lists

@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>.

@prefix s: <http://example.org/students/vocab#>.

<http://example.org/courses/6.001>

s:students

[a rdf:List;

 rdf:first <http://example.org/students/Amy>;

 rdf:rest [a rdf:List

 rdf:first <http://example.org/students/Mohamed>;

 rdf:rest [a rdf:List;

 rdf:first <http://example.org/students/Johann>;

 rdf:rest rdf:nil]]] .

RDF Lists

Turtle has special syntax to represent lists:

```
@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>.
```

```
@prefix s: <http://example.org/students/vocab#>.
```

```
<http://example.org/courses/6.001>
```

```
  s:students (
```

```
    <http://example.org/students/Amy>
```

```
    <http://example.org/students/Mohamed>
```

```
    <http://example.org/students/Johann>
```

```
  ).
```

Critique of RDF

RDF Critique: Properties


- Properties are special kinds of resources
 - Properties can be used as the object in an object-attribute-value triple (statement)
 - Defined **independent** of resources
- This possibility offers flexibility
- But it is unusual for modelling languages and OO programming languages
- It can be confusing for modellers

RDF Critique: Binary Predicates

- RDF uses only binary properties
 - This is a restriction because often we use predicates with more than two arguments
 - But binary predicates can simulate these
- Example: **referee(X, Y, Z)**
 - **X** is the referee in a chess game between players **Y** and **Z**

RDF Critique: Binary Predicates

- We introduce:
 - a new auxiliary resource **chessGame**
 - the binary predicates **ref**, **player1**, and **player2**
- We can represent **referee(X,Y,Z)** as:



RDF Critique: Reification

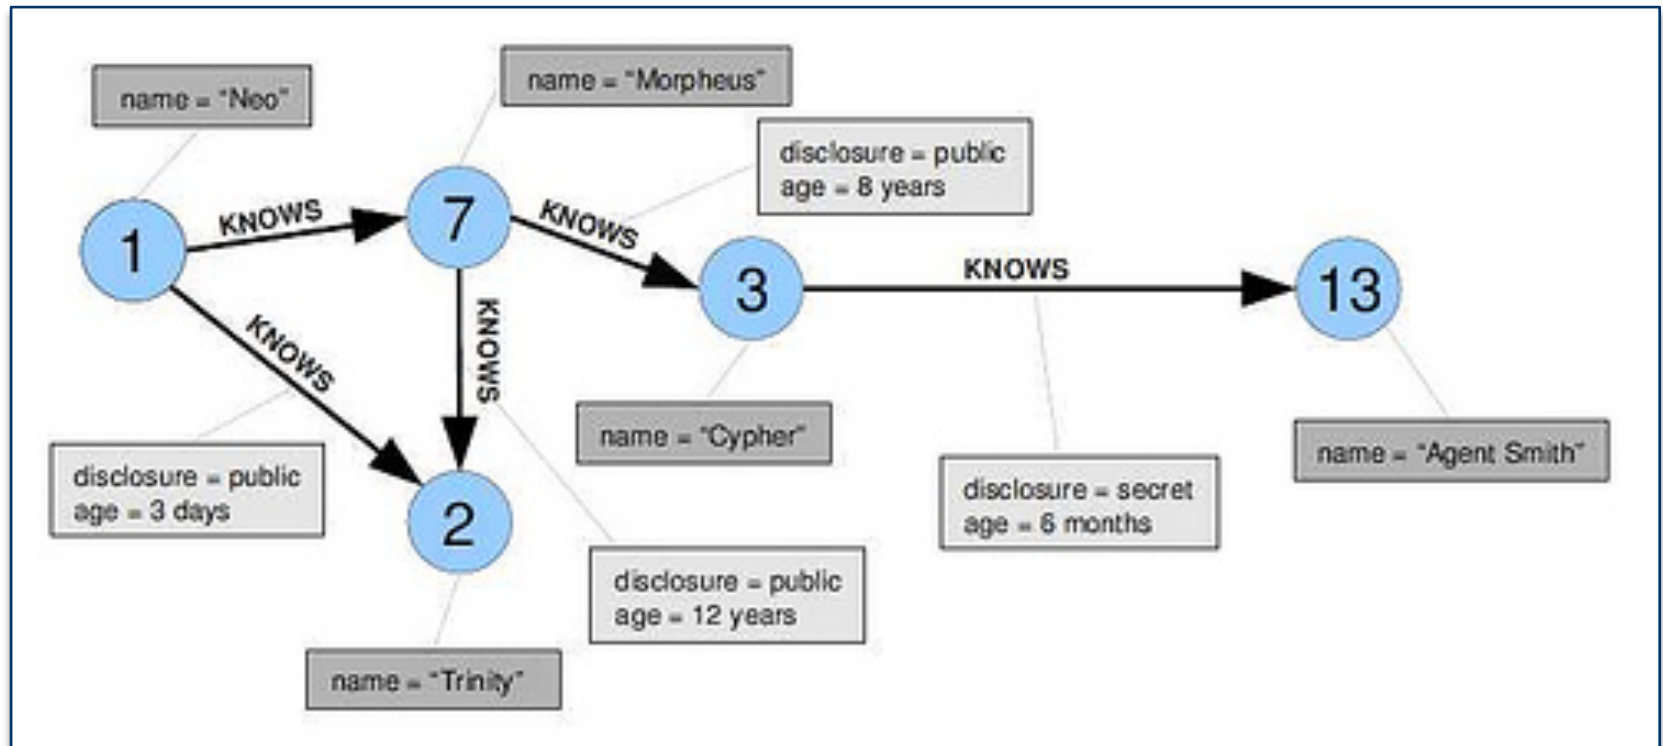
- The reification mechanism is quite powerful
- It appears misplaced in a simple language like RDF
- Making statements about statements introduces a level of complexity that is not necessary for a basic layer of the Semantic Web
- Instead, it would have appeared more natural to include it in more powerful layers, which provide richer representational capabilities

RDF Critique: Graph Representation

- The simple graph or network representation has more drawbacks
- Linear languages introduce ways to represent this with parentheses or a way to represent a block structure
- Scoping, for example, is clumsy at best in RDF
believe(john, and (love(bob, carol), love(carol, bob)))
- Some of these are addressed through the notion of a *named graph* in RDF

RDF graph model is simple

- RDF's graph model is a simple one
- [Neo4J](#) is a popular graph database where both nodes and links can have properties



RDF Critique: Summary

- RDF has its idiosyncrasies and is not an optimal modeling language **but**
- It is already a de facto standard
- It has sufficient expressive power
 - Reasonable foundation on which to build
- Using RDF offers the benefit that information maps unambiguously to a model

Conclusion

Topics

- Basic concepts of RDF
 - Resources, properties, values, statements, triples
 - URIs and URIrefs
 - RDF graphs
 - Literals, qnames
- Vocabularies and modeling
 - Vocabularies
 - Blank nodes, data modeling, types, reification
 - Lists, bags, collections
- Serialization of RDF graphs
 - XML, Turtle, Ntriples
- Critique of RDF