SPARQL

An RDF Query Language

SPARQL

- SPARQL is a recursive acronym for
  SPARQL **P**rotocol **A**nd **R**df **Q**uery **L**anguage
- SPARQL is the SQL for RDF
- Example query suitable for DBpedia
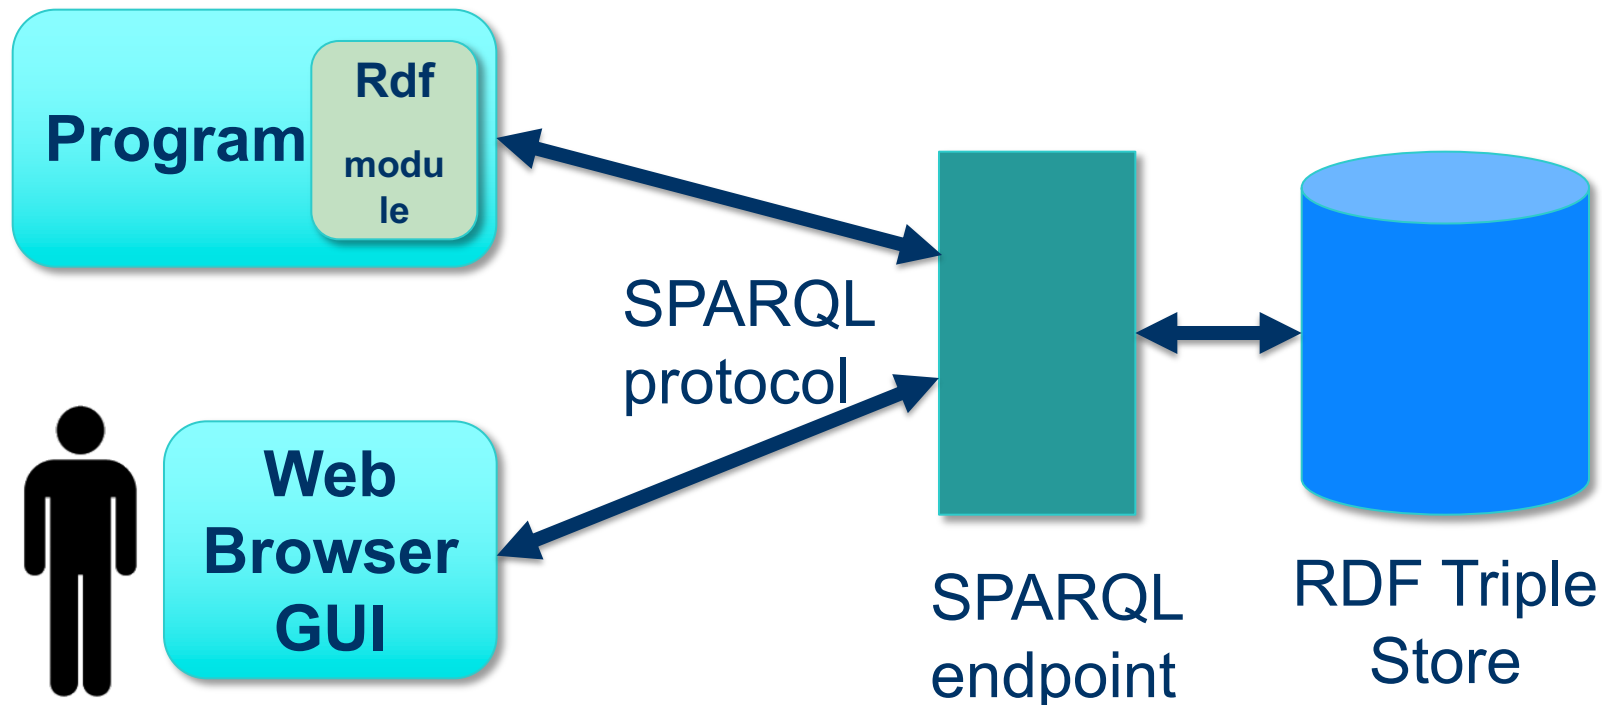
```
 # find countries and their languages
PREFIX dbo: <http://dbpedia.org/ontology/>
SELECT * WHERE {
  ?country a dbo:Country;
          dbo:officialLanguage ?lang .
}
LIMIT 10
```

SPARQL History

- Several RDF query languages were developed prior to SPARQL

- W3C RDF Data Access Working Group (DAWG) worked out SPARQL 2005-2008

- Became a W3C recommendation in Jan 2008

- [SPARQL 1.1](#) (2013) is the current standard

- Support for many prog. languages available

- W3 [SPARQL 1.2](#) Community Group established in 2019 to explore extensions

Typical Architecture

SPARQL endpoint receives queries and requests via HTTP from programs or GUIs, accesses associated RDF triple store and returns result, e.g., data

Program

Rdf module

Web Browser GUI

SPARQL protocol

SPARQL endpoint

RDF Triple Store

Some SPARQL endpoints

There are many public endpoints, e.g.

- Dbpedia: https://dbpedia.ort/sparql/

- Wikidata: https://query.wikidata.org/sparql

- DBLP: https://dblp.l3s.de/d2r/sparql

- See W3C's list of currently alive SPARQL endpoints

It's not hard to set up your own, e.g.

- UMBC cybersecurity knowledge graph: http://eb4.cs.umbc.edu:9090/ckg/query/

Endpoint GUIs

- Some endpoints offer their own SPARQL GUI you can use to enter ad hoc queries

- They may use the same URL as the REST interface and rely on the protocol to know when it's a person and when a query

  - Dbpedia: http://dbpedia.org/sparql/

  - Wikidata: https://query.wikidata.org/

  - DBLP: https://dblp.l3s.de/d2r/snorql/

General SPARQL GUIs

- You can also access or run a general SPARQL GUI that can talk to any SPARQL endpoint

- A nice example is YASGUI, which has a public resource: https://yqagui.org/ and is available to download

- Another open-source GUI is Twinkle

YASGUI: Yet Another SPARQL GUI

https://yasgui.org

SPARQL query structure

- *Prefix declarations* for abbreviating URIs
- *Dataset definition:* what RDF graph(s) are being queried
- *Result clause:* what information to return from the query
- *Query pattern:* what to query for in dataset
- *Query modifiers*, slicing, ordering, rearranging query results

prefix declarations

PREFIX ex: <http://example.com/rdf/> …

optional named graph source

FROM …

result clause (select,ask,update…)

SELECT …

query pattern

WHERE { … }

query modifiers

ORDER BY …

GROUP BY ….

LIMIT 100

Basic SPARQL Query Forms

- SELECT

  Returns all, or a subset of, the variables bound in a query pattern match

- ASK

  Returns boolean indicating whether a query pattern matches or not

- DESCRIBE

  Returns an RDF graph describing resources found

- CONSTRUCT

  Returns an RDF graph constructed by substituting variable bindings in a set of triple templates

A Query: Maryland Cities

```
# find URIs for cities in Maryland
PREFIX yago:  <http://dbpedia.org/class/yago/>
SELECT * WHERE {
    ?city a yago:WikicatCitiesInMaryland
}
```

SPARQL protocol parameters

- To use this query, we need to know]
  - What endpoint (URL) to send it to
  - How we want the results encoded (JSON, XML, …)
  - … other parameters …
- These are set in GUI or your program
  - Except for the endpoint, all have defaults
- Can even query with the unix curl command:

```
curl http://dbpedia.org/sparql/ --data-urlencode query='PREFIX yago:
<http://dbpedia.org/class/yago/> SELECT * WHERE {?city rdf:type
yago:WikicatCitiesInMaryland.}'
```

Maryland Cities and population

```
PREFIX yago: <http://dbpedia.org/class/yago/>t
PREFIX dbo: <http://dbpedia.org/ontology/>
SELECT * WHERE {
    ?city a yago:WikicatCitiesInMaryland;
          dbo:populationTotal ?population .
}
```

Maryland cities, population, names

this returns names in multiple languages ☹
PREFIX yago: <http://dbpedia.org/class/yago/>
PREFIX dbo: <http://dbpedia.org/ontology/>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
SELECT ?city ?name ?population WHERE {
    ?city a yago:WikicatCitiesInMaryland;
            dbo:populationTotal ?population ;
            rdfs:label ?name .
}

Just the @en names, w/o lang tag

```
# FILTER gives conditions that must be true
# LANG(x) returns string's language tag or ""
# STR(x) returns a string's value, i.e. w/o language tag
PREFIX yago:  <http://dbpedia.org/class/yago/>
PREFIX dbo: <http://dbpedia.org/ontology/>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
select (str(?name) as ?name) ?population where {
 ?city a yago:WikicatCitiesInMaryland;
    dbo:populationTotal ?population;
    rdfs:label ?name .
  FILTER (LANG(?name) = "en")
}
```

Order results by population (descending)

sort results by population

PREFIX yago: http://dbpedia.org/class/yago/

PREFIX dbo: <http://dbpedia.org/ontology/>

select str(?name) ?population where {
 ?city a yago:WikicatCitiesInMaryland;
     dbo:populationTotal ?population;
     rdfs:label ?name .
  FILTER (LANG(?name) = "en")
}
ORDER BY DESC(?population)

Wait, where's Catonsville? ☹

- MD's government focused on counties
- Catonsville not considered a city – it has no government
- We need another category of place
    - Census designated place? Populated Place?
- Populated places include counties & regions; let's use census designated place
- But some 'real' cities in Maryland are not listed as census designated places and some are

UNION operator is OR

PREFIX yago: <http://dbpedia.org/class/yago/>

PREFIX dbo: http://dbpedia.org/ontology/

PREFIX dbr: <http://dbpedia.org/resource/>

SELECT str(?name) ?population where {

  {?city dbo:type dbr:Census-designated_place;

       dbo:isPartOf dbr:Maryland .}

  UNION

  {?city a yago:WikicatCitiesInMaryland . }

  ?city dbo:populationTotal ?population; rdfs:label ?name .

  FILTER (LANG(?name) = "en")

}

ORDER BY DESC(?population)

Now we have duplicate entries ☹

- This happens because:
  - Some "cities" are just in WikicatCitiesInMaryland
  - Some are just in Census-designated_places
  - Some are in both
- SPARQL's procedure finds all ways to satisfy a query, and for each one, records the variable bindings
- We add **DISTINCT** to get SPARQL to remove duplicate bindings from the results

DISTINCT produces unique results

PREFIX yago: <http://dbpedia.org/class/yago/>

PREFIX dbo: http://dbpedia.org/ontology/

PREFIX dbr: <http://dbpedia.org/resource/>

SELECT **DISTINCT** str(?name) ?population where {

  {?city dbo:type dbr:Census-designated_place;

       dbo:isPartOf dbr:Maryland .}

  UNION

  {?city a yago:WikicatCitiesInMaryland . }

  ?city dbo:populationTotal ?population;

      rdfs:label ?name .

  FILTER (LANG(?name) = "en")

}

ORDER BY DESC(?population)

Some cities are missing ☹

- Experimentation with query showed there are 427 entities in MD that are either census designated places or cities

- Only get 411 because nine have no population and one has neither a population nor a label

    - Typical of a large and somewhat noisy knowledge graph created from crowdsourced data

- SPARQL's OPIONAL directive to the rescue

OPTIONAL handles missing data

PREFIX yago: <http://dbpedia.org/class/yago/>

PREFIX dbo: <http://dbpedia.org/ontology/>

PREFIX dbr: <http://dbpedia.org/resource/>

select DISTINCT str(?name) ?population where {
  {?city dbo:type dbr:Census-designated_place;
         dbo:isPartOf dbr:Maryland .}
  UNION
  {?city a yago:WikicatCitiesInMaryland . }
  OPTIONAL {?city dbo:populationTotal ?population.}
  OPTIONAL {?city rdfs:label ?name . FILTER (LANG(?name) =
"en") }
}
ORDER BY DESC(?population)

Handling queries with many results

- Endpoints typically have limits on a query's runtime or the number of results it can return

- You can use the LIMIT and OFFSET query modifiers to manage large queries

- Suppose we want to find all types that DBpedia uses

  SELECT distinct ?type WHERE {?x a ?type.}

- DBpedia's public endpoint limits queries to 10K results

Get the first 10K

Get the second 10K with OFFSET

Query 6

```
http://dbpedia.org/sparql
```

```sparql
1  PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
2  PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
3  SELECT distinct ?type WHERE {
4    ?x a ?type .
5  }
6  limit 10000 offset 10000
```

Table | Response | Pivot Table | Google Chart | Geo

Showing 1 to 50 of 10,000 entries (in 20.643 seconds)   Search: [   ]   Show 50 entries

type
1 http://www.wikidata.org/entity/Q2300833
2 http://www.wikidata.org/entity/Q2317783

```python
from SPARQLWrapper import SPARQLWrapper, JSON
default_endpoint = "http://dbpedia.org/sparql"
type_query = """SELECT DISTINCT ?class WHERE {{?x a ?class}} LIMIT {LIM} OFFSET {OFF}"""
def getall(query, endpoint=default_endpoint):
    limit = 10000
    offset = total = 0
    found = limit
    tuples = []
    sparql = SPARQLWrapper(endpoint)
    sparql.setReturnFormat('json')
    while found == limit:  # keep going until we don't get limit results
        q = query.format(LIM=limit, OFF=offset)
        sparql.setQuery(q)
        results = sparql.query().convert()
        found = 0
        for result in results["results"]["bindings"]:
            found += 1
            tuples.append(tuple([str(v['value']) for v in result.values()]))
        print('Found', found, 'results')
        total = total + found
        offset = offset + limit
    return tuples
```

A simple program gets them all

ASK query

- An ASK query returns True if it can be satisfied and False if not

- Was Barack Obama born in the US?

  PREFIX dbo: <http://dbpedia.org/ontology/>

  PREFIX dbr: <http://dbpedia.org/resource/>

  ask WHERE {

   {dbr:Barack_Obama dbo:birthPlace dbr:United_States}

   UNION

   {dbr:Barack_Obama dbo:birthPlace ?x .

    ?x dbo:isPartOf*/dbo:country dbr:United_States }

  }

DESCRIBE Query

- "Describe ?x" means "tell me everything you know about ?x

- Example: Describe Alan Turing …

  DESCRIBE <http://dbpedia.org/resource/Alan_Turing>

  -- or –

  PREFIX dbr: <http://dbpedia.org/resource/>

  DESCRIBE dbr:Alan_Turing

- Returns a collection of ~1500 triples in which dbr:Alan_Turing is either the subject or object

Describes's results?

- The DAWG did not reach a consensus on what describe should return

- Possibilities include
  - All triples where the variable bindings are mentioned
  - All triples where the bindings are the subject
  - Something else

- What is useful might depend on the application or the amount of data involved

- So it was left to the implementation

DESCRIBE Query (2)

- Describe the film "Double Indemnity"

  PREFIX foaf: <http://xmlns.com/foaf/0.1/>

 PREFIX dbo: <http://dbpedia.org/ontology/>

  describe ?x WHERE {

   ?x a dbo:Film; foaf:name ?filmName .

   FILTER (STR(?filmName) = "Double Indemnity")

  }

- Returns a collection of ~500 triples

# DESCRIBE Query (3)

- Describe can return triples about multiple entities

- Describe films directed by Billy Wilder

  PREFIX dbo: http://dbpedia.org/ontology/

  PREFIX dbr: <http://dbpedia.org/resource/>

  describe ?x WHERE {

   ?x a dbo:Film; dbo:director dbr:Billy_Wilder.

  }

- Returns a collection of ~8400 triples about the 27 films he directed

DESCRIBE Query (4)

- Describe can return triples about multiple entities, but you can limit the number

- Describe films directed by Billy Wilder

  PREFIX dbo: http://dbpedia.org/ontology/

  PREFIX dbr: <http://dbpedia.org/resource/>

  describe ?x WHERE {

   ?x a dbo:Film; dbo:director dbr:Billy_Wilder.

   } **LIMIT 1**

- Returns a collection of ~500 triples about just one film, The Apartment.

Construct query (1)

- Construct queries return graphs as results, e.g., film directors and the actors they've directed

    PREFIX dbo: <http://dbpedia.org/ontology/>

    PREFIX ex: <http://example.org/>

    CONSTRUCT {?director ex:directed ?actor}

    WHERE {?film a dbo:Film;

                     dbo:director ?director;

                     dbo:starring ?actor}

- Returns a graph with ~21,000 triples

On construct

- Having a result form that produces an RDF graph is a good idea

- It enables on to construct systems by using the output of one SPARQL query as the data over which another query works

- This kind of capability was a powerful one for relational databases

Construct query (2)

- Construct queries return graphs as results, e.g., film directors and the actors they've directed

    PREFIX dbo: <http://dbpedia.org/ontology/>

    PREFIX ex: <http://example.org/>

    CONSTRUCT {?director ex:directed ?actor}

    WHERE {?film a dbo:Film;

                        dbo:director ?director;

                        dbo:starring ?actor}

- Returns a graph with ~21,000 triples

Construct query (3)

- Actors and directors or producers they've worked for

  PREFIX dbo: <http://dbpedia.org/ontology/>

  PREFIX ex: <http://example.org/>

  Construct {?actor ex:workedFor ?directorOrProducer}

  WHERE {

  ?film a dbo:Film;

  dbo:director|dbo:producer ?directorOrProducer;

  dbo:starring ?actor}

- Returns a graph with ~31,000 triples

> SPARQL 1.1 allows using alternative properties separated by vertical bar

Example: finding missing inverses

- DBpedia is missing many inverse relations, inclueing more than 10k missing spouse relations
- This creates a graph of all the missing ones, which can be added back to the KG via UPDATE ADD

PREFIX dbo: <http://dbpedia.org/ontology/>

CONSTRUCT { ?p2 dbo:spouse ?p1. }

WHERE {?p1 dbo:spouse ?p2.

    FILTER NOT EXISTS {?p2 dbo:spouse ?p1}}

- Not the **NOT EXISTS** operator that succeeds iff its graph pattern is not satisfiable

RDF Named graphs

- Having multiple RDF graphs in a single document/repository and naming them with URIs

- Provides useful additional functionality built on top of the RDF Recommendations

- SPARQL queries can involve several graphs, a background one and multiple named ones, e.g.:

```
SELECT ?who ?g ?mbox
FROM <http://example.org/dft.ttl>
FROM NAMED <http://example.org/alice>
FROM NAMED <http://example.org/bob>
WHERE
{ ?g dc:publisher ?who .
    GRAPH ?g { ?x foaf:mbox ?mbox }
}
```

UPDATE QUERIES

- **Simple insert**

  INSERT DATA { :book1 :title "A new book" ; :creator "A.N.Other" . }

- **Simple delete**

  DELETE DATA { :book1 dc:title "A new book" . }

- Combine the two for a modification, optionally guided by the results of a graph pattern

  PREFIX foaf: <http://xmlns.com/foaf/0.1/>

  DELETE { ?person foaf:givenName 'Bill' }

  INSERT { ?person foaf:givenName 'William' }

  WHERE { ?person foaf:givenName 'Bill' }

Aggregation Operators

- SPARQL 1.1 added many aggregation operators, like count, min, max, …

- Generally used in the results specification

  PREFIX dbo: <http://dbpedia.org/ontology/>

  SELECT **(COUNT(?film) AS ?numberOfFilms)**

  WHERE {?film a dbo:Film .}

- This finds 129,980 films

Group by

- GROUP BY breaks the query's result set into groups before applying the aggregate functions
- Find BO's properties and group them by property and find the number in each group

    PREFIX dbr: <http://dbpedia.org/resource/>

    PREFIX dbo: <http://dbpedia.org/ontology/>

    SELECT ?p (COUNT(?p) as ?number)

    WHERE { dbr:Barack_Obama ?p ?o }

    GROUP BY ?p ORDER BY DESC(count(?p))

Inference via SPARQL

This query adds inverse spouse relations that don't already exist:

> PREFIX dbo: <http://dbpedia.org/ontology/>
>
> INSERT { ?p2 dbo:spouse ?p1. }
>
> WHERE {?p1 dbo:spouse ?p2.
>
> FILTER NOT EXISTS {?p2 dbo:spouse ?p1}}

- SPIN and SHACL are systems to represent simple constraint & inference rules that are done by sparql
- A big feature is that the rules are represented in the graph

SPARQL 1.1 Additions

- SPARQ 1.1 added many more features ...

  - Subqueries

  - Negation: MINUS

  - Federated queries that access multiple endpoints

- Data you want to extract from an RDF graph can probably be returned by one query

  - Might be a complicated one, though ...

- Search web for SPARQL tricks or this book